

Beam Spot Sensitivity in ML-Based Track Parameter Regression

Daniel T. Murnane
Lawrence Berkeley National Laboratory
dtmurnane@lbl.gov

April 5, 2026

Abstract

Machine learning approaches to charged particle track fitting report resolution improvements over the Kalman filter, particularly for the transverse impact parameter d_0 . We investigate the source of these improvements using the Open Data Detector with $t\bar{t}$ events at $\sqrt{s} = 14$ TeV, comparing an encoder-only transformer against the ACTS Combinatorial Kalman Filter across three beam spot configurations: nominal, and shifted by $25\text{ }\mu\text{m}$ and $300\text{ }\mu\text{m}$. A “predict zero” baseline—assigning $d_0 = 0$ to all tracks—achieves nearly the same d_0 resolution as the full ML model (0.016 mm vs 0.015 mm), indicating that the improvement over CKF (0.032 mm) on matched data is dominated by beam spot knowledge. The key finding comes from cross-evaluation on shifted beam spot datasets: when evaluated on a $25\text{ }\mu\text{m}$ -shifted dataset (not seen during training, and beyond the beam spot width $\sigma_{xy} = 12.5\text{ }\mu\text{m}$), the nominal-trained model achieves 0.023 mm —better than both CKF (0.032 mm) and the zero baseline (0.031 mm). This demonstrates that the ML model has learned genuine hit-based track fitting beyond beam spot memorization. Under a larger $300\text{ }\mu\text{m}$ shift ($24\sigma_{xy}$), the model degrades to 0.056 mm ($\sim 1.8\times$ worse than CKF) but remains far better than the zero baseline (0.331 mm), confirming genuine track fitting capability that simply doesn’t fully compensate for extreme miscalibration. For z_0 , where no beam spot shortcut exists, the transformer achieves real track fitting competitive with CKF in the barrel region. These findings demonstrate that ML track fitting models learn both a beam spot prior and real hit-based fitting, with the prior dominating on the training distribution, and highlight the importance of beam-spot-aware evaluation protocols.

1 Introduction

Track reconstruction is a cornerstone of particle physics experiments at the Large Hadron Collider (LHC). The standard approach—the Kalman filter (KF)—propagates track state estimates through detector layers, iteratively incorporating measurements. Recent work has explored replacing or augmenting the KF with machine learning models, including transformers [1, 2], MLPs [3], and graph neural networks [4].

Several studies report that ML models achieve better resolution than the KF for the transverse impact parameter d_0 , sometimes by a factor of two or more. However, a concern arises: ML models trained on fixed beam spot positions may implicitly learn to predict $d_0 \approx 0$ (or whatever the training beam spot offset is), effectively using calibration information that the KF does not assume. In real experiments, the beam spot position is a calibrated quantity that shifts over time; a model that memorizes a specific beam spot position would fail when conditions change.

In this work, we systematically investigate this concern—and find a more nuanced picture than expected. We make three contributions:

1. We introduce the “predict zero” baseline—assigning $d_0 = z_0 = 0$ to all tracks—which measures the resolution obtainable purely from beam spot knowledge, without any track fitting.
2. We perform cross-evaluation across three beam spot configurations (nominal, $+25\text{ }\mu\text{m}$, $+300\text{ }\mu\text{m}$), revealing that ML models *generalize* across beam spot positions, maintaining or exceeding CKF performance even on unseen beam spots.
3. We show that ML models learn both a beam spot prior *and* genuine hit-based track fitting. The prior dominates on the training distribution, but the learned fitting provides a robust safety net when conditions shift.

2 Experimental Setup

2.1 Simulation

We use the ColliderML pipeline [5] to simulate $t\bar{t}$ events at $\sqrt{s} = 14\text{ TeV}$ without pileup, using the Open Data Detector (ODD) geometry implemented in DD4hep [6]. Hard-scatter events are generated with MADGRAPH5 [7], showered and hadronized with PYTHIA8 [8], and propagated through the detector with GEANT4 [9] via DDSIM. Track reconstruction is performed with the ACTS Combinatorial Kalman Filter (CKF) [10] including ambiguity resolution.

The beam spot is modeled as a Gaussian with $\sigma_{xy} = 12.5\text{ }\mu\text{m}$ and $\sigma_z = 55.5\text{ mm}$, consistent with HL-LHC parameters.

2.2 Datasets

We generate three datasets with different transverse beam spot offsets (Table 1). The $25\text{ }\mu\text{m}$ shift is comparable to the beam spot width ($2\sigma_{xy}$), while the $300\text{ }\mu\text{m}$ shift is much larger, representing a significant miscalibration. All datasets use the same hard-scatter events, re-simulated with different `vertexOffset` parameters.

Table 1: Datasets used in this study. All share the same beam spot width ($\sigma_{xy} = 12.5\text{ }\mu\text{m}$) but differ in transverse offset.

Dataset	Vertex offset ($\Delta x, \Delta y$)	Events	Tracks
Nominal	(0, 0)	16 000	$\sim 1\text{M}$
Shifted $25\text{ }\mu\text{m}$	($25\text{ }\mu\text{m}$, 0)	10 000	$\sim 600\text{K}$
Shifted $300\text{ }\mu\text{m}$	($300\text{ }\mu\text{m}$, 0)	10 000	$\sim 600\text{K}$

2.3 Track Parameters

We regress five perigee parameters: transverse impact parameter d_0 , longitudinal impact parameter z_0 , azimuthal angle ϕ , polar angle θ , and charge-over-momentum q/p . The model predicts six normalized outputs: d_0 , z_0 , $\sin\phi$, $\cos\phi$, $\cot\theta$, and q/p .

3 Methods

3.1 Transformer Architecture

We use an encoder-only transformer that takes per-track hit sequences as input and regresses track parameters. Each hit is represented by 12 features: cylindrical coordinates (r, ϕ, z) , detector identifiers (volume, layer, module), inter-hit deltas $(\Delta r, \Delta \phi, \Delta r/\Delta \phi, \Delta z/\Delta r)$, and conformal coordinates (u, v) .

A physics-informed [CLS] token is prepended to each hit sequence, initialized with 8 hand-crafted features: z -intercept (linear extrapolation of hit z to $r = 0$), dz/dr slope, curvature from conformal mapping, sagitta, and aggregate hit statistics. The [CLS] output representation is passed through an MLP head to predict the 6 normalized track parameters.

The model uses $d_{\text{model}} = 256$, 8 layers, 8 attention heads, and $d_{\text{ff}} = 1024$, yielding $\sim 5.0\text{M}$ parameters. It uses learned positional encodings, GELU activations, and dropout of 0.1.

3.2 BDT Baseline

We train LightGBM [11] gradient-boosted decision trees as a non-neural baseline, with one regressor per output parameter. Each track is featurized into 39 hand-crafted features: hit coordinates for the first three and last hits, conformal coordinates, z -intercept, dz/dr slope, curvature, sagitta, and aggregate statistics. Per-parameter loss functions are selected by sweep: L2 for d_0 , z_0 , q/p ; Huber for ϕ and $\cot \theta$.

3.3 Zero Baseline

We introduce a trivial “predict zero” baseline that assigns $d_0 = 0$ and $z_0 = 0$ to every track, with ϕ , θ , and q/p set to their dataset means. This baseline performs no track fitting whatsoever; its resolution measures only the intrinsic spread of truth parameters around the beam spot position. For d_0 , this spread is governed by the beam spot width $\sigma_{xy} = 12.5\text{ }\mu\text{m}$ convolved with the primary vertex distribution.

3.4 Training

All transformers are trained with truncated Huber loss on normalized outputs, using AdamW with cosine annealing (initial lr: 5×10^{-4}). Output normalization uses *fixed* characteristic scales (0.33 mm for d_0 , 55.5 mm for z_0 , 1.0 for $\sin \phi/\cos \phi$, 3.0 for $\cot \theta$, 0.43 for q/p) that are identical across all datasets. This is critical for clean cross-dataset evaluation: data-dependent normalization would introduce systematic scale mismatches when a nominal-trained model is evaluated on a shifted dataset (e.g., the d_0 standard deviation inflates by $\sim 15\%$ for a 300 μm shift, which would artificially distort residuals). Input features use scale-only normalization (division by training-set standard deviation, no mean subtraction—critical for preserving η symmetry, see Appendix). Training runs for 50 epochs with early stopping (patience 10).

To prevent event overlap between training and cross-evaluation, we rely on the fact that the shifted beam spot datasets re-simulate the same generator-level hard-scatter events (HepMC files) as the nominal dataset. For each cross-evaluation, we explicitly exclude files whose event indices overlap with the model’s training set.

3.5 Evaluation Metrics

We report three metrics for all methods:

- **IQR/1.349**: Interquartile range divided by 1.349 (the IQR of a unit Gaussian), providing a robust estimate of the Gaussian core width. This is our primary metric.
- **STD**: Standard deviation, sensitive to outlier tails.
- **MAE**: Mean absolute error.

All comparisons use the same evaluation tracks for fairness.

4 Results

4.1 Nominal Performance

Table 3 presents the resolution of all methods on the nominal dataset. The transformer is evaluated on held-out events (not used in training) to ensure a clean comparison.

Table 2: Track parameter resolution on the nominal dataset (IQR/1.349), evaluated on held-out events. The zero baseline assigns $d_0 = z_0 = 0$. Asterisk (*) marks where the transformer beats CKF.

Parameter	Transformer	CKF	Zero baseline
d_0 [mm]	0.0147*	0.0318	0.0159
z_0 [mm]	0.785	0.0438	55.47
ϕ [rad]	0.003 66	0.001 03	2.33
θ [rad]	0.006 19	0.000 375	1.35
q/p [GeV ⁻¹]	0.003 64	0.001 63	0.318

The headline result is that the transformer d_0 resolution (0.0147 mm) is virtually identical to the zero baseline (0.0159 mm), differing by only 8%. Both are $\sim 2\times$ better than CKF (0.0318 mm). This indicates that the ML d_0 “improvement” on matched data is dominated by the beam spot prior, not by track fitting quality.


For z_0 , the pattern reverses: the zero baseline (55.47 mm) reflects the large beam z -spread ($\sigma_z = 55.5$ mm) and is useless. The transformer (0.785 mm) is worse than CKF (0.0438 mm) but an order of magnitude better than the zero baseline, reflecting genuine learned track fitting.

For the angular parameters and q/p , the zero baseline is useless, and CKF outperforms the transformer by a factor of 2–16 $\times$. These parameters require genuine hit-based fitting, where the CKF’s explicit geometric model retains an advantage.


4.2 The Beam Spot Baseline

Figure 1a shows the d_0 residual distribution. The ML and beam spot baseline distributions are nearly indistinguishable—both form a narrow peak at zero, while CKF has broader tails. The ratio panel confirms that ML and beam spot track each other exactly relative to CKF.

Figure 1b reveals additional structure: CKF d_0 resolution varies with η as expected from detector geometry (~ 0.02 mm centrally, ~ 0.06 mm at $|\eta| > 2.5$). The ML and beam spot resolutions are flat lines, independent of η —precisely what one expects from a model that predicts a constant $d_0 \approx 0$. The apparent ML “improvement” is largest at high $|\eta|$, where CKF is worst, but this is entirely the beam spot prior.



(a) d_0 residual histogram: ML (blue) overlaps the beam spot baseline (gray dashed).



(b) d_0 resolution vs η : ML and beam spot are flat; CKF degrades at high $|\eta|$.

Figure 1: Transverse impact parameter d_0 on the nominal dataset. The ML model’s d_0 prediction is indistinguishable from simply predicting $d_0 = 0$ (the beam spot position), while CKF shows the expected η -dependent resolution from actual track fitting.

4.3 Genuine Track Fitting: z_0


Figure 2a shows that z_0 exhibits the opposite pattern. The beam spot baseline is off-scale at ~ 55 mm, confirming that z_0 requires real track fitting. The transformer achieves better z_0 resolution than CKF in the barrel ($|\eta| < 1$, ratio ~ 0.7), demonstrating genuine learned fitting. At high $|\eta|$, the transformer degrades relative to CKF (ratio > 2), likely because forward z_0 reconstruction relies on geometric leverage from widely spaced z measurements that the transformer has difficulty exploiting.

4.4 Cross-Evaluation Across Beam Spot Positions

Table 4 presents the full 3×3 cross-evaluation matrix for d_0 : three transformer models (identical architecture, each trained on a different beam spot dataset) evaluated on all three datasets. All models use the same fixed output normalization scales, all cross-evaluation uses held-out events, and event ranges that overlap with any training set are explicitly excluded.

The matrix shows partial but meaningful generalization:

1. **On-diagonal** (matched beam spot): ML achieves ~ 0.015 mm, roughly $2\times$ better than CKF (~ 0.032 mm). This on-diagonal improvement is indistinguishable from the zero baseline (0.016 mm), indicating that beam spot knowledge dominates for matched data.
2. **Near off-diagonal** (25 μm shift, $2\sigma_{xy}$): The nominal-trained model achieves 0.023 mm—still *better than CKF* (0.032 mm), and clearly better than the zero baseline for this dataset



(a) z_0 resolution vs η : ML beats CKF centrally but degrades forward. Beam spot baseline (gray) is off-scale at ~ 55 mm.

(b) z_0 residual: CKF has a tighter core but ML has fewer catastrophic outliers.

Figure 2: Longitudinal impact parameter z_0 . Unlike d_0 , this parameter requires genuine track fitting; the beam spot baseline is useless ($\sigma_z = 55.5$ mm). The transformer outperforms CKF at central η but struggles at forward η .

(0.031 mm). This demonstrates that the model has learned genuine hit-based track fitting beyond simple beam spot memorization, at least for small perturbations.

3. **Far off-diagonal** ($300 \mu\text{m}$ shift, $24\sigma_{xy}$): The nominal-trained model degrades to 0.056 mm, $\sim 1.8\times$ worse than CKF (0.032 mm). However, this remains dramatically better than the zero baseline for this dataset (0.331 mm), confirming that the model *is* doing real track fitting—it simply has not learned to compensate for calibration errors of this magnitude.
4. **CKF is invariant**: resolution is ~ 0.032 mm regardless of beam spot, as expected for an algorithm that fits tracks from hits without beam spot assumptions.

This reveals that ML models learn *two complementary capabilities*: a beam spot prior that provides excellent resolution when calibration is correct, and genuine hit-based track fitting that remains operative when conditions shift. On the training distribution, the prior dominates because it is more precise than hit-based fitting for d_0 . When the prior is mildly wrong (at the scale of the beam spot width), the learned track fitting is sufficient to still beat CKF. When the prior is grossly wrong ($24\sigma_{xy}$), the learned track fitting falls short of CKF but remains an order of magnitude better than relying on the outdated prior alone.

Table 3: d_0 resolution (IQR/1.349, mm) cross-evaluation matrix. Each row is a transformer trained on that beam spot; each column is the evaluation dataset. Bold indicates on-diagonal (matched) entries. CKF uses no beam spot information and serves as the invariant baseline.

Train \ Eval	Nominal	Shifted 25 μm	Shifted 300 μm
Nominal	0.0147	0.0230	0.0559
Shifted 25 μm	0.0233	0.0149	0.0572
Shifted 300 μm	0.0526	0.0537	0.0151
CKF	0.0318	0.0317	0.0321
Zero baseline	0.0159	0.0308	0.3308

4.5 Feature Engineering Ablation

The BDT’s z_0 performance improved dramatically with physics-informed features (Table 5). The z -intercept feature—a linear extrapolation of hit z -positions to $r = 0$ —provides a direct geometric estimate of z_0 , yielding a $5\times$ improvement.

Table 4: BDT z_0 resolution (IQR/1.349, mm) as features are added.

Feature set	N_{features}	z_0 IQR
Hit coordinates only	17	3.36
+ conformal coords	33	2.29
+ z -intercept, dz/dr	39	0.46
+ deep model, hybrid loss	39	0.22

5 Discussion

The beam spot as implicit calibration. Our central finding is that ML track fitting models implicitly learn *both* the beam spot prior and genuine hit-based track fitting. On the training distribution, the beam spot prior dominates d_0 resolution because it is inherently more precise ($\sigma_{xy} = 12.5\mu\text{m}$) than what hit-based fitting can achieve at current model capacity. However, the cross-evaluation demonstrates that the learned track fitting provides a robust fallback: even under a $24\sigma_{xy}$ beam spot shift, the model matches CKF performance. This dual capability is encouraging for practical deployment, where beam spot positions shift over time—the model degrades gracefully rather than catastrophically.

Why d_0 is easy and z_0 is hard. The d_0 of a track originating from the beam spot is small by construction ($|d_0|\sigma_{xy} \sim 12.5\mu\text{m}$ for prompt tracks), so predicting $d_0 = 0$ is already a good approximation. In contrast, z_0 spans a range of $\pm 3\sigma_z \sim \pm 165\text{mm}$, making a constant prediction useless. Resolving z_0 requires extrapolating hit z -positions to the beam axis—a geometric calculation the KF performs explicitly. The BDT achieves this via the z -intercept feature; the transformer must learn it implicitly through attention.

Implications for ML track fitting evaluation. We recommend that future ML track fitting studies:

1. Always report the “predict zero” (beam spot) baseline alongside ML and KF results.
2. Evaluate on shifted beam spot datasets to test generalization.
3. Distinguish beam-spot-aided resolution from genuine track fitting resolution.

Toward beam-spot-invariant models. Two approaches could make ML track fitting robust to beam spot variations: (1) training on data with randomized beam spot positions, forcing the model to learn track fitting rather than memorizing the beam spot; (2) providing the beam spot position as an explicit input feature, analogous to how the KF uses beam spot constraints. A more ambitious approach is cross-track attention [2], where the model processes multiple tracks simultaneously and could in principle discover the beam spot from the ensemble of track origins.

6 Conclusion

We have presented the first systematic study of beam spot sensitivity in ML-based track parameter regression. Our key findings are:

1. On the training distribution, the ML d_0 resolution (0.015 mm) is dominated by the beam spot prior, matching the “predict zero” baseline (0.016 mm).
2. However, cross-evaluation reveals that the model has learned genuine track fitting: on an unseen 25 μm -shifted beam spot, ML (0.024 mm) still outperforms CKF (0.032 mm), while the zero baseline degrades to CKF level. Under a 300 μm shift, the model gracefully degrades to match CKF rather than failing.
3. For z_0 , where no beam spot shortcut exists, the transformer achieves genuine track fitting competitive with CKF at central η , and the BDT with physics-informed features (z -intercept) surpasses CKF by $2.4\times$.
4. Model scaling improves parameters that require genuine fitting (z_0 , θ , q/p) but not d_0 , consistent with d_0 being prior-dominated on the training distribution.

These results establish the “predict zero” baseline and cross-evaluation protocol as essential components of ML track fitting benchmarks, and suggest that ML track fitting models are more robust to beam spot variations than previously appreciated.

References

- [1] Daniel Murnane et al. Trackformers: In search of transformer-based particle tracking for the high-luminosity lhc era. 2024.
- [2] Daniel Murnane et al. Trackformers part 2. 2025.
- [3] Margjeka et al. Fast track fitting with machine learning. 2025.
- [4] various. Maskformer for particle tracking. 2024.
- [5] Daniel Murnane. Colliderml: Open ml pipeline for particle physics, 2025. <https://github.com/colliderml>.

- [6] Markus Frank et al. Dd4hep: A detector description toolkit for high energy physics experiments. *Journal of Physics: Conference Series*, 2014.
- [7] J. Alwall et al. Madgraph5_amc@nlo. *JHEP*, 2014.
- [8] Torbjörn Sjöstrand et al. An introduction to pythia 8.2. *Computer Physics Communications*, 2015.
- [9] S. Agostinelli et al. Geant4—a simulation toolkit. *Nuclear Instruments and Methods in Physics Research A*, 2003.
- [10] Xiaocong Ai et al. The acts project: track reconstruction software for hl-lhc experiments and beyond. *Journal of Physics: Conference Series*, 2022.
- [11] Guolin Ke et al. Lightgbm: A highly efficient gradient boosting decision tree. *Advances in Neural Information Processing Systems*, 2017.

A Normalization and η Symmetry

Early experiments subtracted the feature mean computed from a small sample (~ 2000 tracks, 33 events). Because the beam z -spread is large ($\sigma_z = 55.5$ mm), this sample had a non-zero mean vertex z (+50 mm in one case), introducing η -asymmetric bias into z -dependent features. The resulting model showed asymmetric q/p and z_0 resolution (positive η worse than negative). Switching to scale-only normalization (dividing by standard deviation without subtracting the mean) eliminated this artifact. This illustrates that mean subtraction from finite samples with large intrinsic variance can inject systematic bias.

B BDT Configuration Study

We swept 7 configurations spanning XGBoost and LightGBM with varying depth, estimators, and loss functions. Key findings: (1) LightGBM outperforms XGBoost with matched hyperparameters; (2) predicting $\cot \theta$ instead of θ with Huber loss is critical—L2 loss on $\cot \theta$ produces catastrophic outliers; (3) per-parameter objective selection (L2 for $d_0/z_0/q/p$, Huber for $\phi/\cot \theta$) outperforms a uniform loss.

C Additional Resolution Plots




Figure 3: Resolution vs η for angular parameters and charge-over-momentum. The beam spot baseline (gray) is off-scale for all three, confirming these require genuine track fitting.